

```
In [6]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import time
from sklearn.neural_network import MLPClassifier
from utils import plot_weights_distribution
from utils import plot_gradients_distribution

from model import FFNN
```

```
In [7]: def load_mnist():
    print("Loading MNIST dataset")
    X, y = fetch_openml('mnist_784', version=1, return_X_y=True, as_frame=False)

    X = X / 255

    y = y.astype(int)

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=69)

    y_train_onehot = np.zeros((y_train.shape[0], 10))
    y_train_onehot[np.arange(y_train.shape[0]), y_train] = 1

    y_test_onehot = np.zeros((y_test.shape[0], 10))
    y_test_onehot[np.arange(y_test.shape[0]), y_test] = 1

    print(f"Data loaded: X_train: {X_train.shape}, y_train: {y_train.shape}")

    return X_train, X_test, y_train, y_train_onehot, y_test, y_test_onehot
```

```
In [8]: X_train, X_test, y_train, y_train_onehot, y_test, y_test_onehot = load_mnist()
```

```
Loading MNIST dataset
Data loaded: X_train: (56000, 784), y_train: (56000,)
```

```
In [9]: def plot_loss_curves(histories: dict, title: str):
    plt.figure(figsize=(10, 6))
    for label, loss in histories.items():
        plt.plot(loss, label=label)
    plt.title(title)
    plt.xlabel("Epoch")
    plt.ylabel("Training Loss")
    plt.legend()
    plt.grid(True)
    plt.show()

def show_accuracy_table(acc_dict: dict, title="Akurasi Akhir"):
    df = pd.DataFrame.from_dict(acc_dict, orient='index', columns=["Accuracy"])
    display(df.sort_values(by="Accuracy", ascending=False).style.set_caption(title))

def train_and_evaluate(name, model, X_train, y_train, X_test, y_test):
    history = model.train(X_train, y_train, epochs=20, batch_size=32, learning_rate=0.01, verbose=0)
    pred = model.forward(X_test)
    pred_classes = np.argmax(pred, axis=1)
    acc = accuracy_score(np.argmax(y_test, axis=1), pred_classes)
    print(f"{name} Accuracy: {acc:.4f}")
    return history['train_loss'], acc, model
```

```


In [10]: ##### EXPERIMENT 1: WIDTH VARIATION #####
widths = [64, 128, 256]
depth = 1
histories_width = {}
acc_width = {}

for w in widths:
    layers = [(784, None)] + [(w, "relu")] * depth + [(10, "softmax")]
    model = FFNN(layers_config=layers, loss="categorical_crossentropy", wei
ght_init_method="he")
    history, acc, _ = train_and_evaluate(f"Width-{w}", model, X_train, y_tr
ain_onehot, X_test, y_test_onehot)
    histories_width[f"Width-{w}"] = history
    acc_width[f"Width-{w}"] = acc

plot_loss_curves(histories_width, "Loss Curve for Width Variations")
show_accuracy_table(acc_width, "Akurasi Akhir untuk Width Variasi")

```

Width-64 Accuracy: 0.9594  
Width-128 Accuracy: 0.9631  
Width-256 Accuracy: 0.9646



Akurasi Akhir untuk  
Width Variasi

	Accuracy
<b>Width-256</b>	0.964571
<b>Width-128</b>	0.963071
<b>Width-64</b>	0.959357

```

In [11]: ##### EXPERIMENT 2: DEPTH VARIATION #####
depths = [1, 2, 3]
width = 128
histories_depth = {}
acc_depth = {}

for d in depths:
    layers = [(784, None)] + [(width, "relu")] * d + [(10, "softmax")]
    model = FFNN(layers_config=layers, loss="categorical_crossentropy", wei
ght_init_method="he")
    history, acc, _ = train_and_evaluate(f"Depth-{d}", model, X_train, y_tr
ain_onehot, X_test, y_test_onehot)
    histories_depth[f"Depth-{d}"] = history
    acc_depth[f"Depth-{d}"] = acc


plot_loss_curves(histories_depth, "Loss Curve for Depth Variations")
show_accuracy_table(acc_depth, "Akurasi Akhir untuk Depth Variasi")

```

Depth-1 Accuracy: 0.9631

Depth-2 Accuracy: 0.9724

Depth-3 Accuracy: 0.9734



Akurasi Akhir untuk  
Depth Variasi

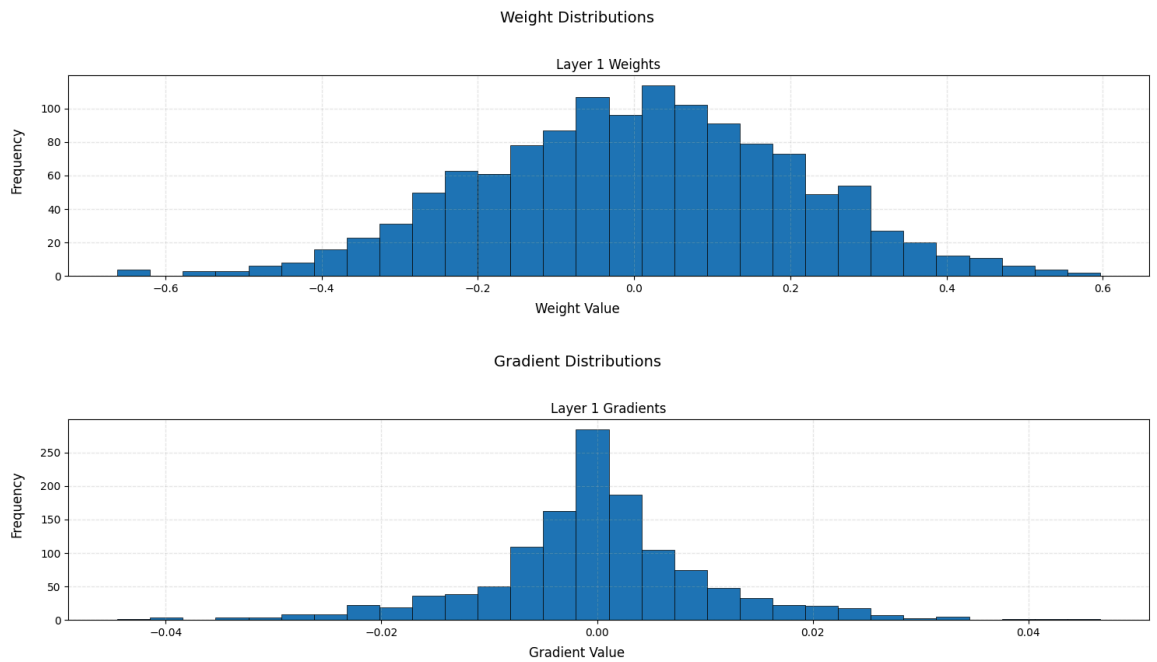
Accuracy	
Depth-3	0.973429
Depth-2	0.972357
Depth-1	0.963071

```
In [12]: ##### EXPERIMENT 3: ACTIVATION FUNCTIONS #####
activations = ['linear', 'relu', 'sigmoid', 'tanh', 'leakyrelu', 'elu']
act_histories = {}
act_acc = {}


for act in activations:
    layers = [(784, None), (128, act), (10, "softmax")]
    model = FFNN(layers_config=layers, loss="categorical_crossentropy", weight_init_method="he")
    history, acc, trained_model = train_and_evaluate(f"Activation-{act}", model, X_train, y_train_onehot, X_test, y_test_onehot)
    act_histories[act] = history
    act_acc[act] = acc
    trained_model.plot_weights_dist([1])
    trained_model.plot_gradients_dist([1])

plot_loss_curves(act_histories, "Loss Curve for Activation Functions")
show_accuracy_table(act_acc, "Akurasi Akhir untuk Aktivasi")
```


Activation-linear Accuracy: 0.9206




Activation-relu Accuracy: 0.9631





Activation-sigmoid Accuracy: 0.9196







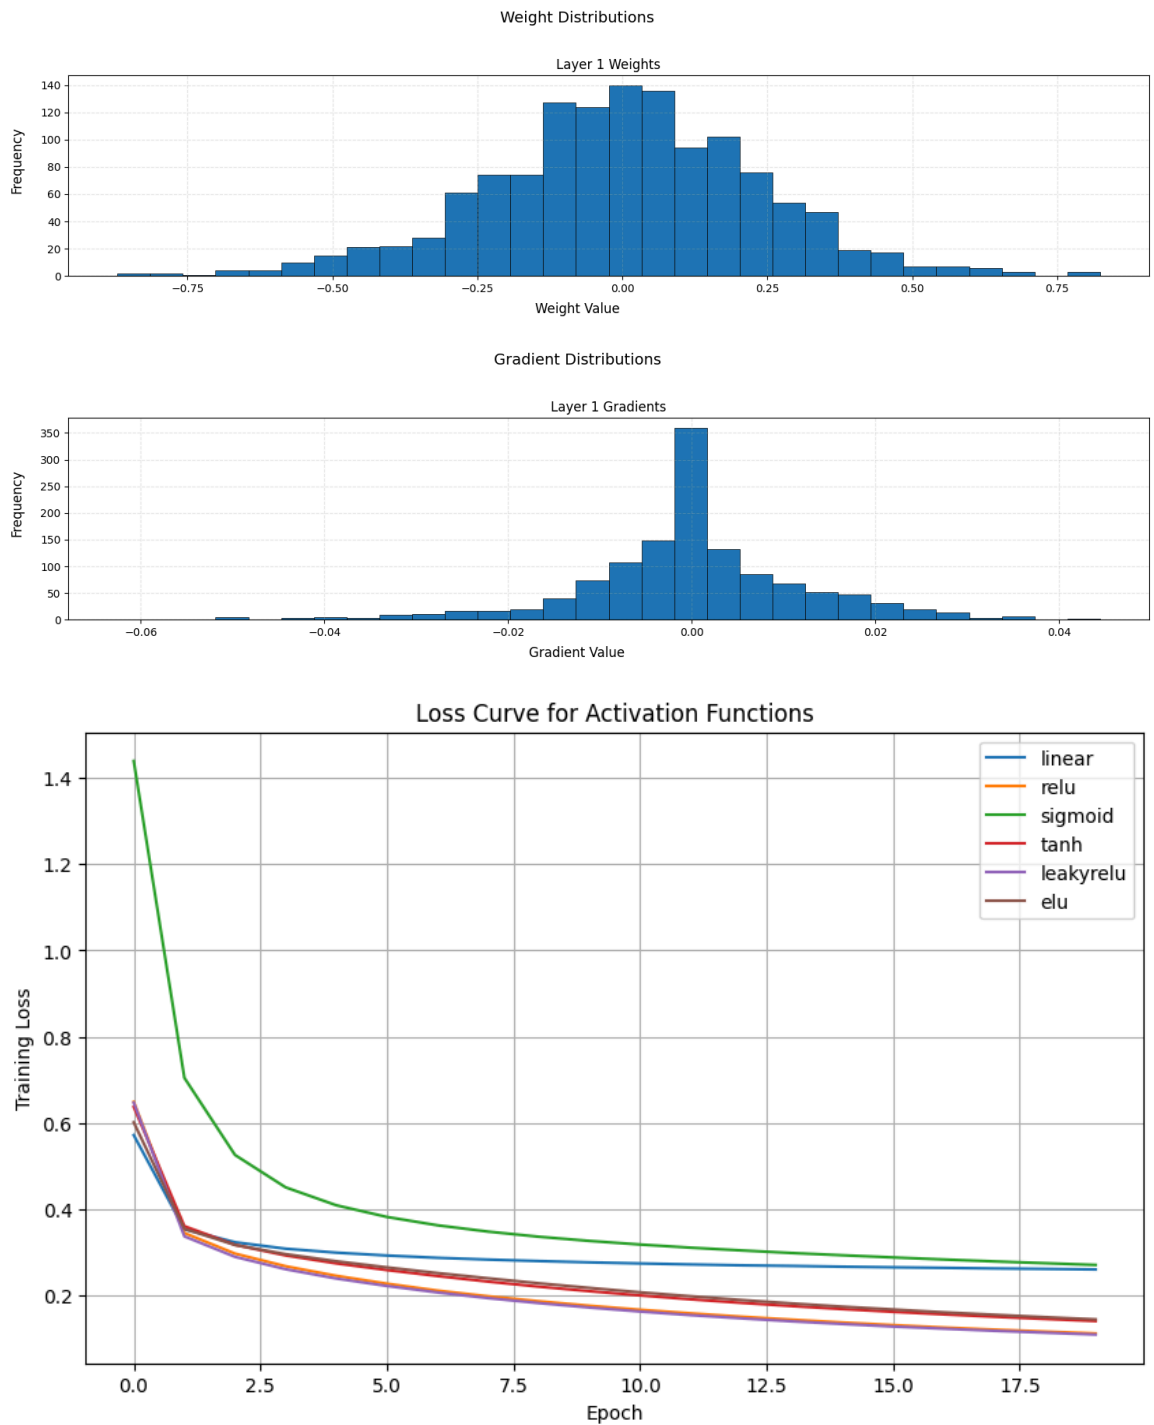
Activation-tanh Accuracy: 0.9539



Activation-leakyrelu Accuracy: 0.9639



Activation-elu Accuracy: 0.9540



Akurasi Akhir untuk  
Aktivasi

	Accuracy
<b>leakyrelu</b>	0.963929
<b>relu</b>	0.963143
<b>elu</b>	0.954000
<b>tanh</b>	0.953929
<b>linear</b>	0.920643
<b>sigmoid</b>	0.919571

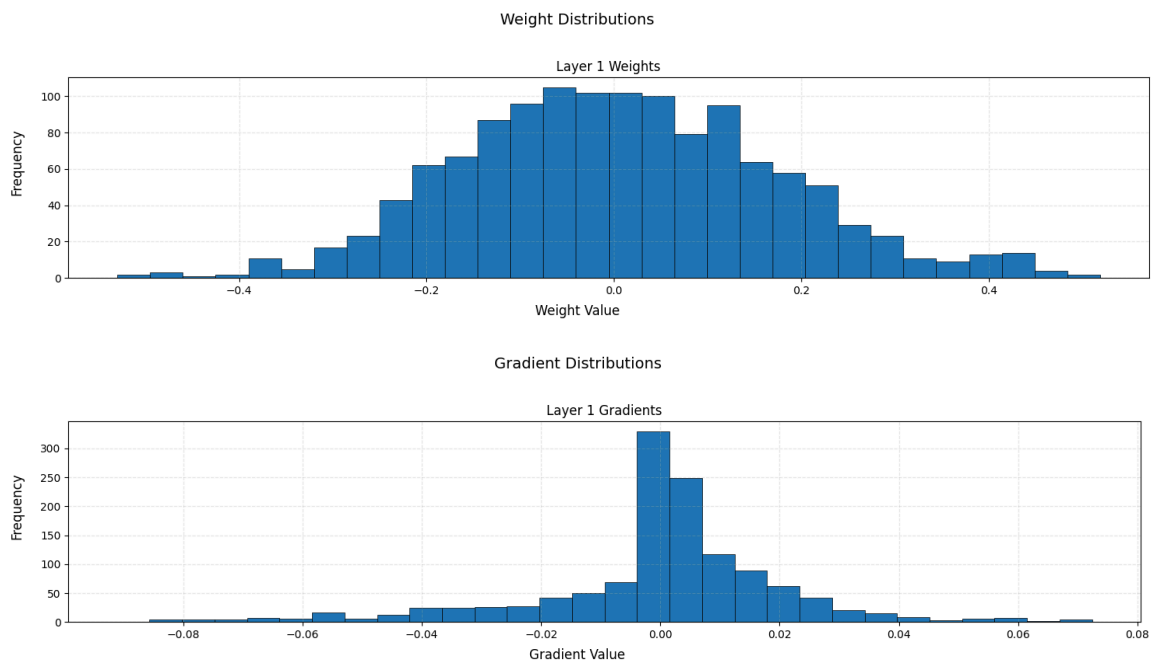


```
In [13]: ##### EXPERIMENT 4: LEARNING RATE #####
lrs = [0.001, 0.01, 0.1]
lr_histories = {}
lr_acc = {}


for lr in lrs:
    layers = [(784, None), (128, "relu"), (10, "softmax")]
    model = FFNN(layers_config=layers, loss="categorical_crossentropy", weight_init_method="he")
    history = model.train(X_train, y_train_onehot, epochs=20, batch_size=32, learning_rate=lr, verbose=0)
    preds = model.forward(X_test)
    acc = accuracy_score(np.argmax(y_test_onehot, axis=1), np.argmax(preds, axis=1))
    print(f"Learning Rate {lr}: {acc:.4f}")
    lr_histories[f"LR-{lr}"] = history['train_loss']
    lr_acc[f"LR-{lr}"] = acc
    model.plot_weights_dist([1])
    model.plot_gradients_dist([1])

plot_loss_curves(lr_histories, "Loss Curve for Learning Rate Variations")
show_accuracy_table(lr_acc, "Akurasi Akhir untuk Learning Rate")
```


Learning Rate 0.001: 0.9106

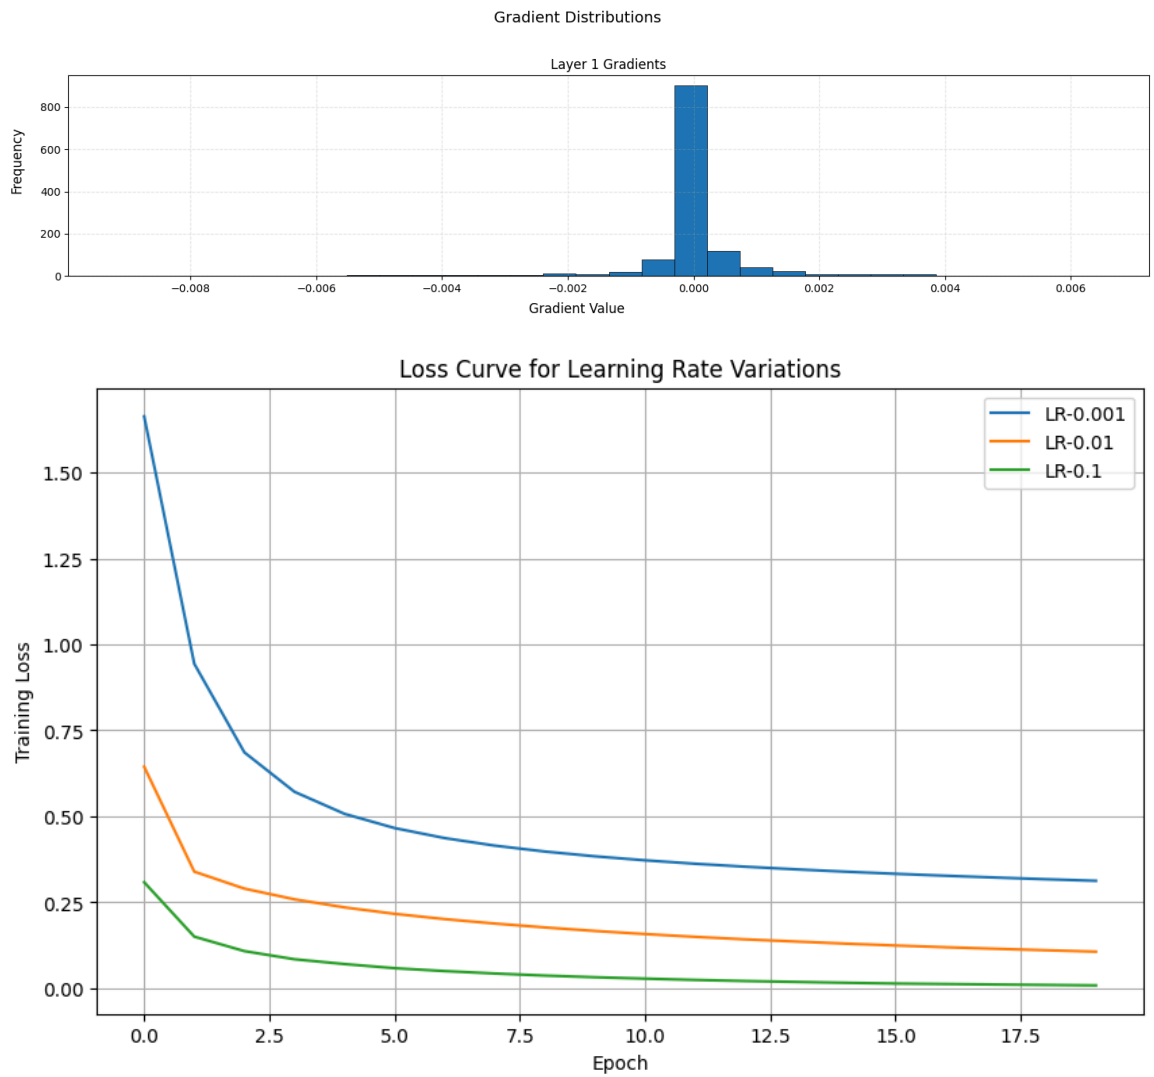


Learning Rate 0.01: 0.9638



Learning Rate 0.1: 0.9793





Akurasi Akhir untuk  
Learning Rate

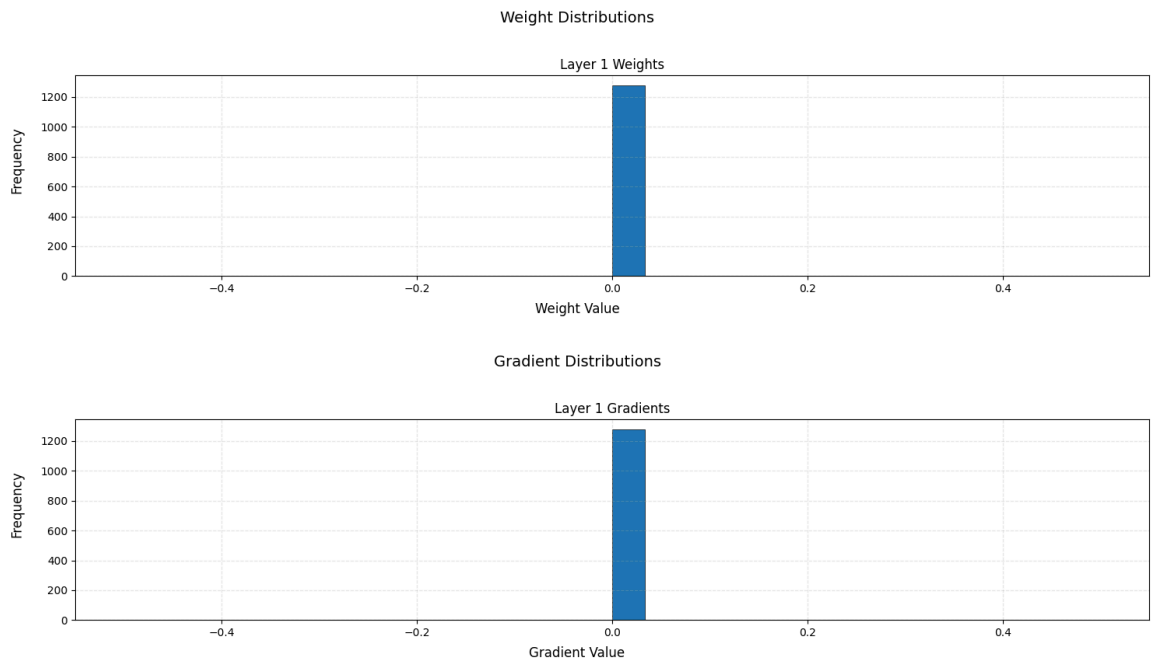
Accuracy	
LR-0.1	0.979286
LR-0.01	0.963786
LR-0.001	0.910571

```
In [14]: ##### EXPERIMENT 5: WEIGHT INITIALIZATION #####
init_methods = ["zero", "random_uniform", "random_normal", "he", "xavier"]
init_histories = {}
init_acc = {}


for init in init_methods:
    model = FFNN(layers_config=[(784, None), (128, "relu"), (10, "softmax")], loss="categorical_crossentropy", weight_init_method=init)
    history, acc, trained_model = train_and_evaluate(f"Init-{init}", model, X_train, y_train_onehot, X_test, y_test_onehot)
    init_histories[init] = history
    init_acc[init] = acc
    trained_model.plot_weights_dist([1])
    trained_model.plot_gradients_dist([1])

plot_loss_curves(init_histories, "Loss Curve for Weight Init Variations")
show_accuracy_table(init_acc, "Akurasi Akhir untuk Weight Init")
```


Init-zero Accuracy: 0.1139




Init-random\_uniform Accuracy: 0.9599





Init-random\_normal Accuracy: 0.9036








Init-he Accuracy: 0.9629



Init-xavier Accuracy: 0.9621





Akurasi Akhir untuk Weight Init

	Accuracy
<b>he</b>	0.962857
<b>xavier</b>	0.962071
<b>random_uniform</b>	0.959857
<b>random_normal</b>	0.903571
<b>zero</b>	0.113857

```
In [15]: ##### EXPERIMENT 6: SKLEARN COMPARISON #####

sklearn_model = MLPClassifier(
    hidden_layer_sizes=(128,), activation='relu', solver='adam', alpha=0.00
01,
    batch_size=32, learning_rate_init=0.01, max_iter=20, random_state=69
)
start = time.time()
sklearn_model.fit(X_train, y_train)
end = time.time()

sklearn_preds = sklearn_model.predict(X_test)
sklearn_acc = accuracy_score(y_test, sklearn_preds)
print(f"sklearn MLPClassifier Accuracy: {sklearn_acc:.4f} (Training Time:
{end - start:.2f}s)")
```

sklearn MLPClassifier Accuracy: 0.9591 (Training Time: 44.02s)

c:\Users\tfpri\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\network\\_multilayer\_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (20) reached and the optimization hasn't converged yet.

```
warnings.warn(
```

```
In [16]: # Custom FFNN dengan hyperparameter yang sama persis dengan sklearn
layers_config = [
    (784, None),
    (128, "relu"),
    (10, "softmax")
]

custom_model = FFNN(
    layers_config=layers_config,
    loss="categorical_crossentropy",
    weight_init_method="he",
    l2_lambda=0.0001
)

start_custom = time.time()
custom_history = custom_model.train(
    X_train, y_train_onehot,
    epochs=20,
    batch_size=32,
    learning_rate=0.01,
    verbose=0
)
end_custom = time.time()

custom_preds = custom_model.forward(X_test)
custom_pred_classes = np.argmax(custom_preds, axis=1)
custom_acc = accuracy_score(y_test, custom_pred_classes)

print("🔍 Hasil Perbandingan Akhir")
print(f"Custom FFNN Accuracy: {custom_acc:.4f} (Waktu: {end_custom - start_custom:.2f}s)")
print(f"Sklearn MLP Accuracy: {sklearn_acc:.4f} (Waktu: {end - start:.2f}s)")
```

🔍 Hasil Perbandingan Akhir  
Custom FFNN Accuracy: 0.9639 (Waktu: 29.91s)  
Sklearn MLP Accuracy: 0.9591 (Waktu: 44.02s)



```
In [17]: configs = {
    "No Regularization": {"l1_lambda": 0.0, "l2_lambda": 0.0},
    "L1 Regularization": {"l1_lambda": 0.001, "l2_lambda": 0.0},
    "L2 Regularization": {"l1_lambda": 0.0, "l2_lambda": 0.001},
    "L1 + L2 Regularization": {"l1_lambda": 0.001, "l2_lambda": 0.001},
}


histories_reg = {}
acc_reg = {}
models_reg = {}

for name, reg_params in configs.items():
    print(f"Training {name}")
    layers_config = [
        (784, None),
        (128, "relu"),
        (10, "softmax")
    ]
    model = FFNN(
        layers_config=layers_config,
        loss="categorical_crossentropy",
        weight_init_method="he",
        l1_lambda=reg_params["l1_lambda"],
        l2_lambda=reg_params["l2_lambda"]
    )
    history, acc, trained_model = train_and_evaluate(
        name, model, X_train, y_train_onehot, X_test, y_test_onehot
    )
    histories_reg[name] = history
    acc_reg[name] = acc
    models_reg[name] = trained_model

plot_loss_curves(histories_reg, "Loss Curve for L1 & L2 Regularization (with combination)")
show_accuracy_table(acc_reg, "Accuracy Comparison for Regularization (with combination)")

for name, model in models_reg.items():
    print(f"\n{name} - Weights Distribution")
    plot_weights_distribution(model, [1])
    print(f"{name} - Gradients Distribution")
    plot_gradients_distribution(model, [1])
```


Training No Regularization  
 No Regularization Accuracy: 0.9619  
 Training L1 Regularization  
 L1 Regularization Accuracy: 0.9622  
 Training L2 Regularization  
 L2 Regularization Accuracy: 0.9634  
 Training L1 + L2 Regularization  
 L1 + L2 Regularization Accuracy: 0.9609




Accuracy Comparison for  
Regularization (with combination)

	Accuracy
<b>L2 Regularization</b>	0.963429
<b>L1 Regularization</b>	0.962214
<b>No Regularization</b>	0.961929
<b>L1 + L2 Regularization</b>	0.960857


No Regularization - Weights Distribution




No Regularization - Gradients Distribution




### L1 Regularization - Weights Distribution




### L1 Regularization - Gradients Distribution




### L2 Regularization - Weights Distribution




### L2 Regularization - Gradients Distribution



### L1 + L2 Regularization - Weights Distribution



## L1 + L2 Regularization - Gradients Distribution



```
In [18]: ## RMSNorm Normalization

configs_rms = {
    "No Normalization": {"rms_norm": False},
    "RMSNorm": {"rms_norm": True}
}

histories_rms = {}
acc_rms = {}
models_rms = {}

for name, rms_params in configs_rms.items():
    print(f"Training {name}")
    layers_config = [
        (784, None),
        (128, "relu", rms_params),
        (10, "softmax")
    ]
    model = FFNN(
        layers_config=layers_config,
        loss="categorical_crossentropy",
        weight_init_method="he"
    )
    history, acc, trained_model = train_and_evaluate(name, model, X_train,
y_train_onehot, X_test, y_test_onehot)
    histories_rms[name] = history
    acc_rms[name] = acc
    models_rms[name] = trained_model

plot_loss_curves(histories_rms, "Loss Curve for RMSNorm Normalization")
show_accuracy_table(acc_rms, "Accuracy Comparison for RMSNorm")


for name, model in models_rms.items():
    print(f"\n{name} - Weights Distribution")
    plot_weights_distribution(model, [1])
    print(f"{name} - Gradients Distribution")
    plot_gradients_distribution(model, [1])
```

Training No Normalization

No Normalization Accuracy: 0.9634

Training RMSNorm

RMSNorm Accuracy: 0.9723



Accuracy Comparison for  
RMSNorm

	Accuracy
<b>RMSNorm</b>	0.972286
<b>No Normalization</b>	0.963429

No Normalization - Weights Distribution